

Parallel Huffman Coding

Francesco Bozzo
DISI, University of Trento
Trento, Italy
francesco.bozzo@studenti.unitn.it
229312

Michele Yin
DISI, University of Trento
Trento, Italy
michele.yin@studenti.unitn.it
229359

Abstract—This report aims to explain the design of a parallel encoding and decoding Huffman algorithm. The application is developed with the C99 programming language, using MPI for multiprocessing and OpenMP for multithreading to scale horizontally with increasing hardware resources. Our tool is both able to exploit multiprocessing to process concurrently multiple files, and to split the Huffman coding of the same file among multiple threads.

Comparing this tool with other online implementations, we can state that...

Index Terms—Huffman, MPI, OpenMP, High Performance Computing

I. INTRODUCTION

Even if the Huffman algorithm is not directly used nowadays, its prefix mechanism is still part of Deflate (PKZIP's algorithm), JPEG, and MP3 compression algorithms

II. THE HUFFMAN ALGORITHM

The Huffman algorithm has the objective of finding a more convenient bit representation to store information through lossless compression. Instead of considering groups of eight bits as the way to encode data, the Huffman algorithm uses variable-length sequences of bits with prefixes defined as *alphabet*.

A. Priority Queue

The Huffman algorithm uses a minimum priority queue to build its alphabet efficiently. This specific data structure ensures logarithmic insertion and deletion time with respect to its size. We implemented the minimum priority queue by using a minimum heap. Practically speaking, to implement the min-heap tree, we used a standard C array ensuring that the min-heap property still holds at every insertion and deletion:

$$A[i] \leq A[l(i)], A[i] \leq A[r(i)] \quad (1)$$

where $A[i]$, $A[l(i)]$, and $A[r(i)]$ are respectively a node, its left child, and its right child in a min-heap tree.

To implement the priority queue also a parallel approach has been considered [1], but since it contains only 256 elements (1 byte), we did not consider it too important in terms of performance.

B. Encoding

The Huffman encoding procedure makes use of a minimum priority queue to build its alphabet efficiently. The idea is to build a tree similar to Figure 1 that defines all the variable-length prefix sequences of bits: these sequences are represented by the path from the tree root to a leaf. Using a greedy approach, the Huffman encoding ensures that the less frequent a byte is in a file, the more probable is to have a longer Huffman representation, which means that its path from the root to its specific leaf is longer.

Once having computed the frequencies for each one of the 256 different bytes in a file, the Huffman algorithm populates a min priority queue with Huffman tree nodes storing the byte value and its frequency. Successively, it removes the least two frequent bytes from the queue, creates a new node with children the two extracted nodes, assigns to it a dummy character and the sum of the frequencies of its two children, and inserts it into the min priority queue. After $n-1$ iterations, the last node is the root of the Huffman tree [2]. Algorithm 1 explains in the detail this procedure.

Algorithm 1: Build the Huffman tree

```
1 // Populate the min priority queue with characters and
  their frequencies
2 for  $i = 1$  to  $n - 1$  do
3   Q.insert(f[i], Tree(f[i], c[i]))
4 // Repeat until the queue has only a single element left
5 for  $i = 1$  to  $n - 1$  do
6   // Get the two least frequent nodes
7   z1 = Q.deleteMin()
8   z2 = Q.deleteMin()
9   // Create an inner tree node and insert it into the
    queue
10  z = Tree(z1.f + z2.f, null)
11  z.left = z1
12  z.right = z2
13  Q.insert(z.f, z)
14 // The last element in the queue is the root of the
    Huffman tree
15 return Q.deleteMin()
```

Once built the Huffman tree, it is possible to generate the




Fig. 1: An example of Huffman tree.

Huffman alphabet visiting the tree using a DFS algorithm, assigning 0 to each left-child traverse and 1 for the right one as presented in Algorithm 2. Finally, it is possible to compress the file by creating a stream of bits corresponding to its content using the Huffman alphabet.

Algorithm 2: Encode using Huffman tree

```

1 while not eof() do
2   bit = read()
3   if bit == 0 then
4     encode(node.left)
5   else
6     encode(node.right)
7   if node is leaf then
8     return node.value

```

C. Decoding

Once having the encoded file and the Huffman tree, it is straightforward to decompress the file to its original shape. The prefix alphabet ensures to have that it is possible to visit the tree using a DFS approach and get a unique decoded version of the file using an approach similar to Algorithm 2.

III. SERIAL HUFFMAN IMPLEMENTATION

To be able to implement the Huffman encoding and decoding algorithm, several things have been taken into account.

- The byte frequencies computed during the encoding process are saved in the compressed file. In this way, the decoding procedure can easily rebuild the Huffman tree.
- Both the encoding and decoding procedures make use of buffers to improve I/O performance: the streams of bytes are first written inside the buffer and then saved on the disk.
- The encoding procedure works with chunks of 4096 bytes. This ensures that the tool can handle even large files when dealing with bit buffers. The decoding procedure deals with chunks of maximum size of $4096 * 32$

bytes, since the maximum length of a Huffman tree traverse from root to leaf is of 256 bits, which is 32 times a byte. As explained later, dealing with chunks makes handy the parallel implementation of the encoding and decoding algorithms.

- In the encoded file, the chunk offsets are saved at the end of the file: this ensures that the compressed stream of prefixed sequences of bits is still meaningful.

IV. PARALLEL HUFFMAN IMPLEMENTATION

To parallelize the serial Huffman algorithm over threads and processes, we decided to follow this simple yet effective structure:

- Multiple processes should handle groups of file and folders separately.
- Multiple threads of the same process should work on different chunks of the same file in parallel.

Specifically, there are multiple reasons behind our parallel architectural design for this project. Here follows a list of the main ones:

- In most operating systems a file is a resource that the OS gives to a single process to avoid race conditions. We wanted to follow a similar design philosophy.
- Most operating systems allow multiple processes to open the same file in reading mode, but only few allow to open the same file in writing mode on multiple processes. This is because that leads to potentially concurrency and data integrity issues. By ensuring to have only a single process that open a specific file, we avoid all these issues.
- Because threads of the same process share the address space, we can avoid the expensive data transfer across processes. When data is read or written to file, there is no need to transfer data between threads.

In the next section, we start by explaining the basic multithreading operations on a single file. Later on, multiprocessing with multiple files is covered.

A. Multithreading

The parallel algorithm for a single file follows the exact same procedure as for the serial version, with just a few key differences to allow multithreading.

Suppose we are dealing with m threads and a single input file. Here follows a common procedure used for both encoding and decoding, with very small differences.

- 1) A single file is divided into chunks of a fixed size. When that is done, the single process forks and creates m threads.
- 2) A single thread (may be the main one, may not) then reads m chunks in a shared memory buffer.
- 3) Each one of the m threads works in parallel on the processing of its assigned chunk using a buffer. Although all the chunks are in shared memory buffer, each thread is assigned only to a single memory section: in this way there is no need to care about racing conditions when writing in the buffer. This is done by creating a `unsigned char buffer[m][buffer_size]`.

- Another detail we noticed on the implementation with two dedicated threads for I/O, is that if there are not enough cores on the CPU or if the operating system decides to allocate all the threads on a single processor, the encoding and decoding times are greatly affected by the scheduler. This is because it might be that the O.S. gives priority to threads that are waiting (for example the writer cannot write any block until the first has finished, even if all others have finished). In the worst case scenario, the multithreading architecture could be even slower than the serial one.

V. PERFORMANCE AND BENCHMARKING

In this section we are going to discuss the performance evaluation of our algorithm. We start from evaluating some other implementations of Huffman algorithm and then focus on the analysis of our own implementation.

We include in this section some graphs and statistics, but all detailed results are included in the appendix at the end of this report.

A. Setup

All the tests have been performed on the HPC2 cluster of the University of Trento. To reduce as much as possible the I/O timings, which are crucial for our application, jobs were run in a very specific setting:

- Threads are allocated in the same node and, if possible, on the same socket to guarantee fast communication between them. This has been achieved by using the MPI option `--map-by`.
- Processes are allocated in different nodes, by using the PBS directive `place=scatter:excl`. This is because in our application different processes do not communicate very often between each other, but they just split their work and use the file-system. We also tested `place=pack:excl` but found worse results.

The provided timings are obtained by averaging the result of three runs with the exact same configurations: this helps to minimize the effect of random variance between runs due to potential hardware congestion. Each job has been submitted to the cluster individually, waiting for the previous to finish preventing hiccups due to multiple instances of the program trying to access the same file-system resource, which we noticed to greatly affect I/O times.

We extensively verified the correctness of our algorithm by encoding and then decoding a file, then comparing the decoded result to the original file by using the `diff` command. Moreover, Valgrind has been used to spot any memory leaks that could cause runtime errors.

B. Datasets

The evaluation dataset has been randomly generated by using a simple C program: we created text files with character frequencies that follow the occurrences of the letters in English language. Although for benchmarking we are only using ASCII characters, please note that our algorithm reads bytes




Fig. 3: Encoding times with and without the flush()

of data and can therefore work with any file. The dataset is composed by files of increasing size: 1 MiB, 5 MiB, 10 MiB, 50 MiB, 100 MiB, 500 MiB, 1 GiB, 5 GiB, and 10 GiB. This should be enough range to understand how different algorithms scale with increasing file sizes.

Secondly, because our algorithm also works with many files at once, we chose some GitHub repository as benchmark: Numpy, PyTorch and Linux kernel as benchmarks. They are respectively 30MiB, 200MiB and 1.3 GiB of total size, with about $\sim 2,000$, $\sim 12,000$, and $\sim 80,000$ files, which on average are each 16MiB in size.

Considering the dataset sizes and amount testing we have done, we estimate to have accessed the disk (both read and writes) for about ~ 6 TiB of data for our whole benchmark analysis.

C. Other works

D. Multithreading evaluation

In this section we evaluate our algorithm on single files, by using multithreading and testing our parallelization performances with an increasing number of threads.

With small files (i.e. less than a few MiBs) the algorithm is actually slower with more threads. This is likely to be the case because the data is not large enough to offset the cost of creating multiple threads. Considering the cost of parallel processing, we highly suggest not doing parallelization for small files.

However, when the data is large enough our multithreaded approach scales up. With the largest file (10 GiB) our algorithm has an efficiency of 93% when tested with 4 threads. This number steadily decreases with the number of threads, and we hit 32% with 24 threads.

Please note that our tool is heavily limited by the cluster I/O, since we can not make it parallel: especially writings to disk require a considerable amount of time. As a proof of this, if we consider only the processing section of our algorithm, and ignore the time required by the `fflush()` operation before a `fclose()`, we see significantly decreased times. Although we start from similar efficiency at low threading, we achieve a 54% efficiency with 24 threads in this case: if we consider only CPU time, which does not count the time spent waiting in I/O, we get very promising results.

We also noticed worse times in the decoding procedure with respect to encoding. This is probably due to the fact that the




Fig. 4: Encoding vs Decoding times




Fig. 5: Barrier vs Lock based synchronization performance

tool is reading fixed chunks of 4096 bytes in encoding, but during decoding the chunk can have a variable length.

As already discussed previously, we also tested a version on a lock-based synchronization instead of barriers. Although theoretically we should get better results, in practice we obtain almost equal performances across the board, with only the largest files having some significant difference. This difference is greater in decoding than in encoding.

E. Multiprocessing

We tested on different folders as benchmark the parallelization performances of our algorithm. Theoretically, our algorithm should scale very well with increasing number of processes and files, because once the main process sends the list of files to encode to each process there is no further communication between different processes. However, we find that real world results are different. While we don't expect much increasing in the performance by the number of threads, because the files are on average very small, we should get better results by using many processes. In reality our testing shows that this is not the case. Our efficiencies are extremely low. We believe the main limitation of our algorithm is the I/O of the cluster. When tested on a local machine (MacBook

A2442) with one thread and 1,2,4 processes, we find that our algorithm does indeed scale with the number of processes. We also find that `scatter` scales better than `pack`. Probably because with `scatter` we are not limited by the I/O of a single node.

F. Memory, storage and other considerations

Our memory requirements don't scale upwards with the size of the files we need to encode, because at any given time, the most we only store is the size of two buffers, each is a `unsigned char buffer[num_threads][4096]`. These buffers are emptied at each I/O cycle, therefore leading to a high memory efficiency. Only when processing multiple files our memory requirements scale up with the number of files, as each process needs to store information about the files is assigned to process.

We also find that we achieve on average 48% compression rate of our synthetic data. Instead, on Linux kernel, because contains more varied data which consists of mostly code, we achieve a lower 62% compression rate. Testing on other already compressed data as `.zip` or `AV1` resulted in a compression rate of 100%.

VI. CONCLUSIONS

To conclude.

REFERENCES

- [1] G. S. Brodal, J. L. Träff, and C. D. Zaroliagis, "A parallel priority queue with constant time operations," *Journal of Parallel and Distributed Computing*, vol. 49, no. 1, pp. 4–21, 1998. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0743731598914253>
- [2] A. Bertossi and A. Montresor, *Algoritmi e strutture di dati*. CittàStudi, 2010. [Online]. Available: <https://books.google.es/books?id=WKWL5QAACA>




Fig. 6: Encoding times

VII. APPENDIX

A. Encoding

TABLE I: Overall encoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|---------|---------|--------|---------|--------|---------------|--------------|---------------|--------------|
| 1 MiB | 0.108 | 0.103 | 0.083 | 0.091 | 0.072 | 0.110 | 0.134 | 0.211 | 0.170 | 0.030 | 0.025 | 0.027 | 0.026 |
| 5 MiB | 0.407 | 0.201 | 0.126 | 0.104 | 0.097 | 0.079 | 0.079 | 0.077 | 0.091 | 0.181 | 0.056 | 0.058 | 0.053 |
| 10 MiB | 0.761 | 0.474 | 0.266 | 0.208 | 0.180 | 0.186 | 0.220 | 0.335 | 0.243 | 0.102 | 0.089 | 0.096 | 0.170 |
| 50 MiB | 3.924 | 2.138 | 1.279 | 1.020 | 0.857 | 0.759 | 0.777 | 0.658 | 0.752 | 0.424 | 0.537 | 0.417 | 0.394 |
| 100 MiB | 7.727 | 4.068 | 2.403 | 1.860 | 1.554 | 1.432 | 1.299 | 1.247 | 1.231 | 0.856 | 0.945 | 0.720 | 0.888 |
| 500 MiB | 35.775 | 18.383 | 11.496 | 8.351 | 7.559 | 5.755 | 6.086 | 5.570 | 5.143 | 3.800 | 3.669 | 5.378 | 3.538 |
| 1 GiB | 73.112 | 38.168 | 22.507 | 16.336 | 14.213 | 13.756 | 12.620 | 10.826 | 10.335 | 7.810 | 7.468 | 10.505 | 8.655 |
| 5 GiB | 359.849 | 186.629 | 99.804 | 77.462 | 65.655 | 67.306 | 52.909 | 48.459 | 48.013 | 36.637 | 38.346 | 59.600 | 47.238 |
| 10 GiB | 694.975 | 370.549 | 188.228 | 136.172 | 113.588 | 104.025 | 93.723 | 116.825 | 83.208 | 76.289 | 73.011 | 65.807 | 85.032 |

TABLE II: Pure encoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|--------|--------|--------|--------|--------|--------------|---------------|---------------|--------------|
| 1 MiB | 0.064 | 0.034 | 0.018 | 0.013 | 0.010 | 0.009 | 0.039 | 0.094 | 0.035 | 0.004 | 0.004 | 0.004 | 0.004 |
| 5 MiB | 0.298 | 0.152 | 0.078 | 0.054 | 0.041 | 0.034 | 0.030 | 0.024 | 0.021 | 0.018 | 0.017 | 0.015 | 0.013 |
| 10 MiB | 0.655 | 0.352 | 0.184 | 0.125 | 0.098 | 0.081 | 0.073 | 0.153 | 0.057 | 0.034 | 0.029 | 0.031 | 0.040 |
| 50 MiB | 3.303 | 1.755 | 0.915 | 0.624 | 0.488 | 0.433 | 0.343 | 0.266 | 0.295 | 0.177 | 0.139 | 0.178 | 0.132 |
| 100 MiB | 6.721 | 3.446 | 1.827 | 1.245 | 0.990 | 0.815 | 0.687 | 0.542 | 0.486 | 0.380 | 0.482 | 0.242 | 0.462 |
| 500 MiB | 32.383 | 16.605 | 8.682 | 6.246 | 4.862 | 4.059 | 3.413 | 2.662 | 2.385 | 1.929 | 1.386 | 3.187 | 1.233 |
| 1 GiB | 67.176 | 34.097 | 17.735 | 12.655 | 9.879 | 8.237 | 7.004 | 5.431 | 4.892 | 3.982 | 2.936 | 5.939 | 4.277 |
| 5 GiB | 331.253 | 174.089 | 90.562 | 63.491 | 48.645 | 38.946 | 33.185 | 26.100 | 23.623 | 19.312 | 14.208 | 44.206 | 32.709 |
| 10 GiB | 639.337 | 328.580 | 170.409 | 119.682 | 95.123 | 79.805 | 68.875 | 53.342 | 47.573 | 39.444 | 40.491 | 29.375 | 56.942 |

TABLE III: Overall encoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.527 | 0.328 | 0.200 | 0.189 | 0.099 | 0.068 | 0.032 | 0.032 | 0.151 | 0.135 | 0.083 | 0.064 |
| 5 MiB | 1.000 | 1.014 | 0.811 | 0.654 | 0.525 | 0.518 | 0.430 | 0.332 | 0.225 | 0.094 | 0.229 | 0.146 | 0.119 |
| 10 MiB | 1.000 | 0.803 | 0.716 | 0.609 | 0.527 | 0.409 | 0.288 | 0.142 | 0.156 | 0.311 | 0.267 | 0.165 | 0.070 |
| 50 MiB | 1.000 | 0.918 | 0.767 | 0.641 | 0.572 | 0.517 | 0.421 | 0.373 | 0.261 | 0.386 | 0.229 | 0.196 | 0.156 |
| 100 MiB | 1.000 | 0.950 | 0.804 | 0.693 | 0.621 | 0.539 | 0.496 | 0.387 | 0.314 | 0.376 | 0.255 | 0.224 | 0.136 |
| 500 MiB | 1.000 | 0.973 | 0.778 | 0.714 | 0.592 | 0.622 | 0.490 | 0.401 | 0.348 | 0.392 | 0.305 | 0.139 | 0.158 |
| 1 GiB | 1.000 | 0.958 | 0.812 | 0.746 | 0.643 | 0.531 | 0.483 | 0.422 | 0.354 | 0.390 | 0.306 | 0.145 | 0.132 |
| 5 GiB | 1.000 | 0.964 | 0.901 | 0.774 | 0.685 | 0.535 | 0.567 | 0.464 | 0.375 | 0.409 | 0.293 | 0.126 | 0.119 |
| 10 GiB | 1.000 | 0.938 | 0.923 | 0.851 | 0.765 | 0.668 | 0.618 | 0.372 | 0.418 | 0.380 | 0.297 | 0.220 | 0.128 |

TABLE IV: Pure encoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.943 | 0.868 | 0.830 | 0.796 | 0.747 | 0.137 | 0.042 | 0.091 | 0.626 | 0.495 | 0.326 | 0.233 |
| 5 MiB | 1.000 | 0.979 | 0.952 | 0.925 | 0.902 | 0.873 | 0.840 | 0.765 | 0.719 | 0.676 | 0.540 | 0.422 | 0.346 |
| 10 MiB | 1.000 | 0.929 | 0.887 | 0.870 | 0.837 | 0.806 | 0.748 | 0.267 | 0.574 | 0.810 | 0.713 | 0.441 | 0.254 |
| 50 MiB | 1.000 | 0.941 | 0.903 | 0.882 | 0.845 | 0.764 | 0.802 | 0.777 | 0.559 | 0.779 | 0.741 | 0.386 | 0.392 |
| 100 MiB | 1.000 | 0.975 | 0.920 | 0.900 | 0.848 | 0.824 | 0.815 | 0.776 | 0.692 | 0.736 | 0.436 | 0.578 | 0.227 |
| 500 MiB | 1.000 | 0.975 | 0.932 | 0.864 | 0.832 | 0.798 | 0.791 | 0.760 | 0.679 | 0.699 | 0.730 | 0.212 | 0.410 |
| 1 GiB | 1.000 | 0.985 | 0.947 | 0.885 | 0.850 | 0.816 | 0.799 | 0.773 | 0.687 | 0.703 | 0.715 | 0.236 | 0.245 |
| 5 GiB | 1.000 | 0.951 | 0.914 | 0.870 | 0.851 | 0.851 | 0.832 | 0.793 | 0.701 | 0.715 | 0.729 | 0.156 | 0.158 |
| 10 GiB | 1.000 | 0.973 | 0.938 | 0.890 | 0.840 | 0.801 | 0.774 | 0.749 | 0.672 | 0.675 | 0.493 | 0.453 | 0.175 |

B. Encoding with Locks

TABLE V: Overall encoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|---------|---------|--------|--------|--------|--------|--------------|---------------|--------------|
| 1 MiB | 0.088 | 0.084 | 0.073 | 0.065 | 0.072 | 0.072 | 0.067 | 0.067 | 0.075 | 0.071 | 0.031 | 0.027 | 0.033 |
| 5 MiB | 0.425 | 0.406 | 0.275 | 0.172 | 0.167 | 0.164 | 0.148 | 0.139 | 0.139 | 0.136 | 0.058 | 0.064 | 0.059 |
| 10 MiB | 0.713 | 0.678 | 0.242 | 0.194 | 0.176 | 0.147 | 0.168 | 0.129 | 0.140 | 0.157 | 0.102 | 0.094 | 0.102 |
| 50 MiB | 3.585 | 3.371 | 1.141 | 0.838 | 0.748 | 0.698 | 0.961 | 0.968 | 0.733 | 0.895 | 0.430 | 0.388 | 0.573 |
| 100 MiB | 6.805 | 6.601 | 2.159 | 1.622 | 1.359 | 1.200 | 1.098 | 0.974 | 0.888 | 0.842 | 0.771 | 0.722 | 0.735 |
| 500 MiB | 33.464 | 31.995 | 10.564 | 7.854 | 6.683 | 6.041 | 5.518 | 5.931 | 4.577 | 4.248 | 4.140 | 3.551 | 3.811 |
| 1 GiB | 68.489 | 65.942 | 21.173 | 16.425 | 13.648 | 12.077 | 10.966 | 9.795 | 9.036 | 8.342 | 8.007 | 8.001 | 7.690 |
| 5 GiB | 367.228 | 370.850 | 110.299 | 84.128 | 75.319 | 66.408 | 59.666 | 59.479 | 51.181 | 48.904 | 38.824 | 38.677 | 40.676 |
| 10 GiB | 676.529 | 642.922 | 187.620 | 135.991 | 123.822 | 109.770 | 96.643 | 93.226 | 90.825 | 88.154 | 63.886 | 61.847 | 64.598 |

TABLE VI: Pure encoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|--------|--------|--------|--------|--------|--------|--------------|--------------|---------------|
| 1 MiB | 0.061 | 0.062 | 0.036 | 0.028 | 0.030 | 0.022 | 0.024 | 0.024 | 0.030 | 0.016 | 0.011 | 0.009 | 0.011 |
| 5 MiB | 0.361 | 0.363 | 0.209 | 0.109 | 0.103 | 0.092 | 0.079 | 0.066 | 0.053 | 0.050 | 0.023 | 0.023 | 0.017 |
| 10 MiB | 0.617 | 0.607 | 0.169 | 0.117 | 0.090 | 0.076 | 0.078 | 0.054 | 0.053 | 0.059 | 0.039 | 0.030 | 0.035 |
| 50 MiB | 3.131 | 3.057 | 0.847 | 0.595 | 0.464 | 0.416 | 0.671 | 0.652 | 0.435 | 0.500 | 0.186 | 0.145 | 0.160 |
| 100 MiB | 6.079 | 6.037 | 1.664 | 1.151 | 0.891 | 0.735 | 0.642 | 0.521 | 0.438 | 0.386 | 0.327 | 0.285 | 0.249 |
| 500 MiB | 31.001 | 30.471 | 8.287 | 5.810 | 4.519 | 3.778 | 3.226 | 2.610 | 2.227 | 1.946 | 1.814 | 1.430 | 1.404 |
| 1 GiB | 62.748 | 62.355 | 17.230 | 11.939 | 9.153 | 7.582 | 6.505 | 5.363 | 4.612 | 4.101 | 3.698 | 2.966 | 2.833 |
| 5 GiB | 343.876 | 351.067 | 99.909 | 73.825 | 62.694 | 54.050 | 47.849 | 42.794 | 33.944 | 29.952 | 18.749 | 14.847 | 14.376 |
| 10 GiB | 629.449 | 623.710 | 172.829 | 120.552 | 94.623 | 78.728 | 68.025 | 55.215 | 46.723 | 41.474 | 32.493 | 27.176 | 23.935 |

TABLE VII: Overall encoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.526 | 0.300 | 0.224 | 0.152 | 0.122 | 0.109 | 0.082 | 0.059 | 0.051 | 0.089 | 0.067 | 0.042 |
| 5 MiB | 1.000 | 0.523 | 0.386 | 0.410 | 0.317 | 0.259 | 0.239 | 0.191 | 0.153 | 0.130 | 0.229 | 0.138 | 0.113 |
| 10 MiB | 1.000 | 0.525 | 0.737 | 0.613 | 0.506 | 0.485 | 0.354 | 0.346 | 0.254 | 0.189 | 0.219 | 0.158 | 0.109 |
| 50 MiB | 1.000 | 0.532 | 0.786 | 0.713 | 0.599 | 0.513 | 0.311 | 0.231 | 0.245 | 0.167 | 0.261 | 0.192 | 0.098 |
| 100 MiB | 1.000 | 0.515 | 0.788 | 0.699 | 0.626 | 0.567 | 0.516 | 0.437 | 0.383 | 0.337 | 0.276 | 0.196 | 0.145 |
| 500 MiB | 1.000 | 0.523 | 0.792 | 0.710 | 0.626 | 0.554 | 0.505 | 0.353 | 0.366 | 0.328 | 0.253 | 0.196 | 0.137 |
| 1 GiB | 1.000 | 0.519 | 0.809 | 0.695 | 0.627 | 0.567 | 0.520 | 0.437 | 0.379 | 0.342 | 0.267 | 0.178 | 0.139 |
| 5 GiB | 1.000 | 0.495 | 0.832 | 0.728 | 0.609 | 0.553 | 0.513 | 0.386 | 0.359 | 0.313 | 0.296 | 0.198 | 0.141 |
| 10 GiB | 1.000 | 0.526 | 0.901 | 0.829 | 0.683 | 0.616 | 0.583 | 0.454 | 0.372 | 0.320 | 0.331 | 0.228 | 0.164 |

TABLE VIII: Pure encoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.493 | 0.425 | 0.367 | 0.252 | 0.283 | 0.209 | 0.160 | 0.101 | 0.161 | 0.172 | 0.138 | 0.085 |
| 5 MiB | 1.000 | 0.497 | 0.431 | 0.551 | 0.438 | 0.391 | 0.380 | 0.344 | 0.341 | 0.300 | 0.491 | 0.331 | 0.334 |
| 10 MiB | 1.000 | 0.508 | 0.913 | 0.881 | 0.852 | 0.816 | 0.655 | 0.714 | 0.583 | 0.436 | 0.495 | 0.426 | 0.275 |
| 50 MiB | 1.000 | 0.512 | 0.924 | 0.878 | 0.843 | 0.753 | 0.389 | 0.300 | 0.360 | 0.261 | 0.525 | 0.449 | 0.306 |
| 100 MiB | 1.000 | 0.503 | 0.914 | 0.880 | 0.853 | 0.827 | 0.789 | 0.729 | 0.693 | 0.656 | 0.581 | 0.444 | 0.381 |
| 500 MiB | 1.000 | 0.509 | 0.935 | 0.889 | 0.857 | 0.821 | 0.801 | 0.742 | 0.696 | 0.664 | 0.534 | 0.452 | 0.345 |
| 1 GiB | 1.000 | 0.503 | 0.910 | 0.876 | 0.857 | 0.828 | 0.804 | 0.731 | 0.680 | 0.638 | 0.530 | 0.441 | 0.346 |
| 5 GiB | 1.000 | 0.490 | 0.860 | 0.776 | 0.686 | 0.636 | 0.599 | 0.502 | 0.507 | 0.478 | 0.573 | 0.483 | 0.374 |
| 10 GiB | 1.000 | 0.505 | 0.911 | 0.870 | 0.832 | 0.800 | 0.771 | 0.712 | 0.674 | 0.632 | 0.605 | 0.483 | 0.411 |

TABLE IX: Overall decoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|---------|---------|---------|---------------|---------|---------|---------|----------------|--------------|
| 1 MiB | 0.130 | 0.091 | 0.079 | 0.099 | 0.091 | 0.097 | 0.118 | 0.115 | 0.119 | 0.104 | 0.045 | 0.026 | 0.030 |
| 5 MiB | 0.422 | 0.253 | 0.177 | 0.166 | 0.141 | 0.156 | 0.151 | 0.139 | 0.130 | 0.121 | 0.087 | 0.087 | 0.082 |
| 10 MiB | 0.871 | 0.517 | 0.384 | 0.296 | 0.293 | 0.433 | 0.372 | 0.356 | 0.368 | 0.355 | 0.171 | 0.129 | 0.161 |
| 50 MiB | 4.108 | 2.277 | 1.511 | 1.252 | 1.276 | 1.382 | 1.412 | 1.409 | 1.434 | 1.328 | 0.693 | 0.594 | 0.544 |
| 100 MiB | 7.961 | 4.507 | 2.915 | 2.275 | 1.884 | 2.418 | 2.400 | 2.422 | 2.343 | 2.251 | 4.392 | 1.858 | 1.722 |
| 500 MiB | 37.619 | 22.511 | 13.501 | 11.334 | 10.035 | 9.326 | 9.640 | 11.182 | 11.069 | 10.021 | 19.667 | 5.680 | 7.006 |
| 1 GiB | 76.679 | 43.252 | 27.585 | 23.271 | 18.590 | 17.824 | 18.532 | 22.027 | 19.813 | 17.128 | 37.634 | 12.183 | 14.964 |
| 5 GiB | 361.871 | 193.112 | 108.812 | 96.790 | 81.066 | 84.897 | 79.142 | 73.725 | 79.374 | 75.976 | 155.945 | 111.347 | 81.739 |
| 10 GiB | 759.072 | 391.202 | 212.381 | 156.396 | 135.844 | 135.401 | 134.098 | 138.000 | 134.006 | 132.255 | 253.313 | 106.227 | 119.210 |

TABLE X: Pure decoding times

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|---------|---------|---------|---------|---------|--------|--------|---------|--------|--------|---------|---------------|--------------|
| 1 MiB | 0.083 | 0.058 | 0.052 | 0.074 | 0.069 | 0.071 | 0.092 | 0.090 | 0.092 | 0.076 | 0.015 | 0.005 | 0.005 |
| 5 MiB | 0.340 | 0.188 | 0.123 | 0.101 | 0.084 | 0.094 | 0.089 | 0.079 | 0.073 | 0.062 | 0.020 | 0.018 | 0.016 |
| 10 MiB | 0.752 | 0.414 | 0.267 | 0.191 | 0.191 | 0.265 | 0.267 | 0.246 | 0.282 | 0.249 | 0.059 | 0.032 | 0.029 |
| 50 MiB | 3.634 | 1.888 | 1.058 | 0.745 | 0.670 | 0.976 | 0.954 | 0.943 | 0.966 | 0.857 | 0.271 | 0.148 | 0.132 |
| 100 MiB | 7.132 | 3.613 | 2.012 | 1.388 | 1.123 | 1.540 | 1.520 | 1.541 | 1.478 | 1.468 | 2.433 | 0.299 | 0.511 |
| 500 MiB | 36.395 | 18.594 | 9.874 | 6.980 | 5.458 | 4.707 | 5.739 | 6.970 | 6.499 | 5.473 | 16.056 | 1.514 | 2.716 |
| 1 GiB | 74.764 | 38.408 | 19.764 | 13.840 | 10.908 | 10.242 | 9.231 | 12.504 | 12.060 | 9.825 | 29.059 | 3.072 | 5.671 |
| 5 GiB | 359.883 | 188.886 | 97.788 | 82.355 | 56.999 | 59.936 | 50.887 | 39.291 | 48.562 | 48.877 | 141.437 | 15.227 | 55.262 |
| 10 GiB | 754.842 | 385.021 | 199.850 | 138.739 | 113.105 | 95.262 | 81.112 | 108.727 | 94.317 | 81.262 | 245.681 | 32.584 | 75.318 |

TABLE XI: Overall decoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.719 | 0.411 | 0.220 | 0.179 | 0.134 | 0.092 | 0.071 | 0.055 | 0.052 | 0.090 | 0.103 | 0.067 |
| 5 MiB | 1.000 | 0.835 | 0.595 | 0.423 | 0.374 | 0.270 | 0.233 | 0.189 | 0.163 | 0.145 | 0.152 | 0.101 | 0.080 |
| 10 MiB | 1.000 | 0.842 | 0.567 | 0.491 | 0.371 | 0.201 | 0.195 | 0.153 | 0.118 | 0.102 | 0.159 | 0.140 | 0.085 |
| 50 MiB | 1.000 | 0.902 | 0.680 | 0.547 | 0.402 | 0.297 | 0.242 | 0.182 | 0.143 | 0.129 | 0.185 | 0.144 | 0.118 |
| 100 MiB | 1.000 | 0.883 | 0.683 | 0.583 | 0.528 | 0.329 | 0.276 | 0.205 | 0.170 | 0.147 | 0.057 | 0.089 | 0.072 |
| 500 MiB | 1.000 | 0.836 | 0.697 | 0.553 | 0.469 | 0.403 | 0.325 | 0.210 | 0.170 | 0.156 | 0.060 | 0.138 | 0.084 |
| 1 GiB | 1.000 | 0.886 | 0.695 | 0.549 | 0.516 | 0.430 | 0.345 | 0.218 | 0.194 | 0.187 | 0.064 | 0.131 | 0.080 |
| 5 GiB | 1.000 | 0.937 | 0.831 | 0.623 | 0.558 | 0.426 | 0.381 | 0.307 | 0.228 | 0.198 | 0.073 | 0.068 | 0.069 |
| 10 GiB | 1.000 | 0.970 | 0.894 | 0.809 | 0.698 | 0.561 | 0.472 | 0.344 | 0.283 | 0.239 | 0.094 | 0.149 | 0.099 |

TABLE XII: Pure decoding efficiency

| Threads Size | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|-----------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.720 | 0.401 | 0.188 | 0.150 | 0.117 | 0.075 | 0.058 | 0.045 | 0.046 | 0.170 | 0.338 | 0.245 |
| 5 MiB | 1.000 | 0.902 | 0.690 | 0.558 | 0.504 | 0.361 | 0.317 | 0.269 | 0.233 | 0.227 | 0.519 | 0.404 | 0.324 |
| 10 MiB | 1.000 | 0.909 | 0.705 | 0.655 | 0.493 | 0.284 | 0.234 | 0.191 | 0.133 | 0.126 | 0.397 | 0.490 | 0.402 |
| 50 MiB | 1.000 | 0.962 | 0.859 | 0.813 | 0.678 | 0.372 | 0.317 | 0.241 | 0.188 | 0.177 | 0.419 | 0.511 | 0.429 |
| 100 MiB | 1.000 | 0.987 | 0.886 | 0.856 | 0.794 | 0.463 | 0.391 | 0.289 | 0.241 | 0.202 | 0.092 | 0.497 | 0.218 |
| 500 MiB | 1.000 | 0.979 | 0.922 | 0.869 | 0.834 | 0.773 | 0.528 | 0.326 | 0.280 | 0.277 | 0.071 | 0.501 | 0.209 |
| 1 GiB | 1.000 | 0.973 | 0.946 | 0.900 | 0.857 | 0.730 | 0.675 | 0.374 | 0.310 | 0.317 | 0.080 | 0.507 | 0.206 |
| 5 GiB | 1.000 | 0.953 | 0.920 | 0.728 | 0.789 | 0.600 | 0.589 | 0.572 | 0.371 | 0.307 | 0.080 | 0.492 | 0.102 |
| 10 GiB | 1.000 | 0.980 | 0.944 | 0.907 | 0.834 | 0.792 | 0.776 | 0.434 | 0.400 | 0.387 | 0.096 | 0.483 | 0.157 |

TABLE XIII: Overall decoding times

| Threads Sizes | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|------------------|---------|---------|---------|---------|---------|---------|---------|---------|---------|----------------|---------|--------------|---------------|
| 1 MiB | 0.272 | 0.212 | 0.116 | 0.101 | 0.070 | 0.068 | 0.066 | 0.069 | 0.075 | 0.093 | 0.095 | 0.036 | 0.036 |
| 5 MiB | 0.635 | 0.915 | 1.265 | 0.655 | 0.684 | 0.289 | 0.482 | 0.298 | 0.330 | 0.180 | 0.562 | 0.124 | 0.254 |
| 10 MiB | 0.861 | 0.501 | 0.333 | 0.271 | 0.269 | 0.258 | 0.236 | 0.217 | 0.190 | 0.183 | 0.604 | 0.440 | 0.404 |
| 50 MiB | 3.991 | 2.248 | 1.462 | 1.176 | 1.160 | 1.088 | 1.017 | 0.913 | 0.837 | 0.775 | 1.409 | 0.761 | 0.726 |
| 100 MiB | 7.615 | 33.645 | 15.250 | 11.628 | 13.018 | 10.761 | 9.192 | 3.661 | 5.691 | 4.475 | 2.175 | 1.312 | 1.464 |
| 500 MiB | 37.082 | 20.422 | 13.236 | 11.344 | 10.774 | 10.441 | 9.767 | 8.661 | 7.387 | 7.051 | 10.581 | 9.984 | 6.591 |
| 1 GiB | 73.855 | 41.185 | 26.929 | 23.051 | 17.977 | 18.781 | 19.731 | 17.837 | 16.171 | 15.437 | 19.766 | 14.432 | 12.636 |
| 5 GiB | 352.386 | 475.087 | 265.894 | 204.661 | 156.227 | 139.894 | 127.649 | 109.245 | 100.918 | 89.890 | 104.173 | 65.082 | 63.837 |
| 10 GiB | 714.907 | 376.029 | 202.712 | 190.071 | 435.224 | 145.673 | 133.319 | 122.774 | 117.529 | 115.262 | 158.281 | 150.057 | 148.834 |

TABLE XIV: Pure decoding times

| Threads Sizes | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|------------------|---------|---------|---------|----------|---------|---------|---------|--------|--------|--------|---------|---------------|---------------|
| 1 MiB | 0.116 | 0.064 | 0.056 | 0.056 | 0.041 | 0.041 | 0.038 | 0.041 | 0.038 | 0.035 | 0.061 | 0.006 | 0.007 |
| 5 MiB | 0.565 | 0.859 | 1.202 | 0.598 | 0.627 | 0.234 | 0.425 | 0.244 | 0.269 | 0.113 | 0.019 | 0.017 | 0.016 |
| 10 MiB | 0.715 | 0.399 | 0.234 | 0.174 | 0.169 | 0.159 | 0.136 | 0.119 | 0.095 | 0.085 | 0.132 | 0.031 | 0.028 |
| 50 MiB | 3.531 | 1.848 | 1.034 | 0.742 | 0.730 | 0.656 | 0.620 | 0.486 | 0.417 | 0.352 | 0.703 | 0.143 | 0.117 |
| 100 MiB | 6.874 | 33.366 | 14.884 | 10.786 | 12.377 | 9.907 | 8.348 | 2.879 | 5.046 | 3.695 | 0.968 | 0.285 | 0.510 |
| 500 MiB | 35.450 | 18.238 | 10.009 | 7.179 | 6.792 | 6.317 | 5.622 | 4.667 | 3.831 | 3.280 | 6.218 | 1.394 | 1.212 |
| 1 GiB | 72.114 | 36.926 | 20.113 | 14.503 | 11.600 | 12.692 | 11.302 | 9.281 | 7.654 | 6.478 | 9.440 | 2.841 | 2.705 |
| 5 GiB | 349.890 | 471.340 | 257.714 | 197.061 | 140.154 | 124.853 | 110.634 | 91.401 | 74.582 | 62.539 | 73.525 | 14.510 | 12.708 |
| 10 GiB | 711.985 | 370.520 | 191.792 | 182.3175 | 421.918 | 126.678 | 113.482 | 92.082 | 76.309 | 64.483 | 100.849 | 29.389 | 54.947 |

TABLE XV: Overall decoding efficiency

| Threads Sizes | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|------------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.642 | 0.585 | 0.450 | 0.486 | 0.402 | 0.346 | 0.245 | 0.182 | 0.122 | 0.089 | 0.157 | 0.119 |
| 5 MiB | 1.000 | 0.347 | 0.126 | 0.162 | 0.116 | 0.220 | 0.110 | 0.133 | 0.096 | 0.147 | 0.035 | 0.107 | 0.039 |
| 10 MiB | 1.000 | 0.860 | 0.647 | 0.529 | 0.400 | 0.334 | 0.303 | 0.248 | 0.227 | 0.196 | 0.045 | 0.041 | 0.033 |
| 50 MiB | 1.000 | 0.888 | 0.683 | 0.566 | 0.430 | 0.367 | 0.327 | 0.273 | 0.238 | 0.215 | 0.089 | 0.109 | 0.086 |
| 100 MiB | 1.000 | 0.113 | 0.125 | 0.109 | 0.073 | 0.071 | 0.069 | 0.130 | 0.067 | 0.071 | 0.109 | 0.121 | 0.081 |
| 500 MiB | 1.000 | 0.908 | 0.700 | 0.545 | 0.430 | 0.355 | 0.316 | 0.268 | 0.251 | 0.219 | 0.110 | 0.077 | 0.088 |
| 1 GiB | 1.000 | 0.897 | 0.686 | 0.534 | 0.514 | 0.393 | 0.312 | 0.259 | 0.228 | 0.199 | 0.117 | 0.107 | 0.091 |
| 5 GiB | 1.000 | 0.371 | 0.331 | 0.287 | 0.282 | 0.252 | 0.230 | 0.202 | 0.175 | 0.163 | 0.106 | 0.113 | 0.086 |
| 10 GiB | 1.000 | 0.951 | 0.881 | 0.625 | 0.205 | 0.491 | 0.447 | 0.364 | 0.304 | 0.258 | 0.141 | 0.099 | 0.075 |

TABLE XVI: Pure decoding efficiency

| Threads Sizes | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 16 | 20 | 24 | 32 | 48 | 64 |
|------------------|-------|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 MiB | 1.000 | 0.901 | 0.515 | 0.349 | 0.357 | 0.285 | 0.255 | 0.176 | 0.152 | 0.138 | 0.060 | 0.428 | 0.260 |
| 5 MiB | 1.000 | 0.329 | 0.118 | 0.157 | 0.113 | 0.242 | 0.111 | 0.145 | 0.105 | 0.208 | 0.918 | 0.703 | 0.551 |
| 10 MiB | 1.000 | 0.897 | 0.764 | 0.684 | 0.529 | 0.449 | 0.437 | 0.376 | 0.377 | 0.351 | 0.169 | 0.486 | 0.402 |
| 50 MiB | 1.000 | 0.955 | 0.854 | 0.793 | 0.605 | 0.538 | 0.474 | 0.454 | 0.423 | 0.418 | 0.157 | 0.516 | 0.471 |
| 100 MiB | 1.000 | 0.103 | 0.115 | 0.106 | 0.069 | 0.069 | 0.069 | 0.149 | 0.068 | 0.078 | 0.222 | 0.502 | 0.210 |
| 500 MiB | 1.000 | 0.972 | 0.885 | 0.823 | 0.652 | 0.561 | 0.525 | 0.475 | 0.463 | 0.450 | 0.178 | 0.530 | 0.457 |
| 1 GiB | 1.000 | 0.976 | 0.896 | 0.829 | 0.777 | 0.568 | 0.532 | 0.486 | 0.471 | 0.464 | 0.239 | 0.529 | 0.417 |
| 5 GiB | 1.000 | 0.371 | 0.339 | 0.296 | 0.312 | 0.280 | 0.264 | 0.239 | 0.235 | 0.233 | 0.149 | 0.502 | 0.430 |
| 10 GiB | 1.000 | 0.9606 | 0.928 | 0.651 | 0.211 | 0.562 | 0.523 | 0.483 | 0.467 | 0.460 | 0.221 | 0.505 | 0.202 |

E. numpy encoding pack

TABLE XVII: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------------|---------------|---------------|
| 1 | 39.668 | 39.316 | 38.853 |
| 2 | 28.366 | 29.230 | 29.015 |
| 4 | 25.048 | 25.980 | 24.888 |
| 8 | 26.048 | 25.678 | 25.688 |
| 10 | 25.551 | 25.874 | 25.738 |
| 12 | 26.608 | 26.560 | 24.980 |
| 16 | 26.316 | 24.455 | 46.843 |

TABLE XVIII: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.504 | 0.255 |
| 2 | 0.699 | 0.339 | 0.171 |
| 4 | 0.396 | 0.191 | 0.100 |
| 8 | 0.190 | 0.097 | 0.048 |
| 10 | 0.155 | 0.077 | 0.039 |
| 12 | 0.124 | 0.062 | 0.033 |
| 16 | 0.094 | 0.051 | 0.013 |

F. numpy encoding scatter

TABLE XIX: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------------|---------------|---------------|
| 1 | 39.346 | 42.755 | 39.929 |
| 2 | 28.056 | 27.240 | 26.568 |
| 4 | 20.548 | 19.755 | 33.501 |
| 8 | 16.588 | 15.468 | 15.198 |
| 10 | 27.227 | 21.108 | 20.999 |
| 12 | 26.551 | 25.964 | 25.659 |
| 16 | 20.122 | 19.575 | 19.687 |

TABLE XX: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.460 | 0.246 |
| 2 | 0.701 | 0.361 | 0.185 |
| 4 | 0.479 | 0.249 | 0.073 |
| 8 | 0.296 | 0.159 | 0.081 |
| 10 | 0.145 | 0.093 | 0.047 |
| 12 | 0.123 | 0.063 | 0.032 |
| 16 | 0.122 | 0.063 | 0.031 |

G. numpy decoding pack

TABLE XXI: Overall decoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------------|---------------|---------------|
| 1 | 43.866 | 42.586 | 43.507 |
| 2 | 31.064 | 33.149 | 36.228 |
| 4 | 34.853 | 33.162 | 30.350 |
| 8 | 30.040 | 27.924 | 27.418 |
| 10 | 28.016 | 28.269 | 27.062 |
| 12 | 28.995 | 29.148 | 29.472 |
| 16 | 30.624 | 30.456 | 32.091 |

TABLE XXII: Overall decoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.515 | 0.252 |
| 2 | 0.706 | 0.331 | 0.151 |
| 4 | 0.315 | 0.165 | 0.090 |
| 8 | 0.183 | 0.098 | 0.050 |
| 10 | 0.157 | 0.078 | 0.041 |
| 12 | 0.126 | 0.063 | 0.031 |
| 16 | 0.090 | 0.045 | 0.021 |

H. numpy decoding scatter

TABLE XXIII: Overall decoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|--------|---------------|---------------|
| 1 | 46.248 | 44.085 | 43.422 |
| 2 | 30.331 | 28.251 | 28.600 |
| 4 | 22.234 | 21.866 | 21.471 |
| 8 | 18.648 | 16.757 | 16.952 |
| 10 | 16.364 | 15.773 | 15.839 |
| 12 | 15.574 | 15.720 | 15.148 |
| 16 | 22.756 | 22.066 | 22.733 |

TABLE XXIV: Overall decoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.525 | 0.266 |
| 2 | 0.762 | 0.409 | 0.202 |
| 4 | 0.520 | 0.264 | 0.135 |
| 8 | 0.310 | 0.172 | 0.085 |
| 10 | 0.283 | 0.147 | 0.073 |
| 12 | 0.247 | 0.123 | 0.064 |
| 16 | 0.127 | 0.065 | 0.032 |

I. pytorch encoding pack

TABLE XXV: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------|----------------|----------------|
| 1 | 247.430 | 229.251 | 225.370 |
| 2 | 174.185 | 168.884 | 168.979 |
| 4 | 157.589 | 155.282 | 157.478 |
| 8 | 161.628 | 158.789 | 156.443 |
| 10 | 188.020 | 167.726 | 158.378 |
| 12 | 161.222 | 158.014 | 154.675 |
| 16 | 159.562 | 153.257 | 148.363 |

TABLE XXVI: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.540 | 0.274 |
| 2 | 0.710 | 0.366 | 0.183 |
| 4 | 0.393 | 0.199 | 0.098 |
| 8 | 0.191 | 0.097 | 0.049 |
| 10 | 0.132 | 0.074 | 0.039 |
| 12 | 0.128 | 0.065 | 0.033 |
| 16 | 0.097 | 0.050 | 0.026 |

J. pytorch encoding scatter

TABLE XXVII: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------|----------------|----------------|
| 1 | 317.663 | 268.323 | 265.387 |
| 2 | 184.717 | 179.515 | 173.750 |
| 4 | 139.132 | 123.383 | 126.569 |
| 8 | 105.950 | 102.119 | 100.966 |
| 10 | 98.206 | 94.839 | 91.242 |
| 12 | 95.794 | 96.229 | 90.712 |
| 16 | 87.427 | 84.979 | 84.262 |

TABLE XXVIII: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.592 | 0.299 |
| 2 | 0.860 | 0.442 | 0.229 |
| 4 | 0.571 | 0.322 | 0.157 |
| 8 | 0.375 | 0.194 | 0.098 |
| 10 | 0.323 | 0.167 | 0.087 |
| 12 | 0.276 | 0.138 | 0.073 |
| 16 | 0.227 | 0.117 | 0.059 |

K. pytorch encoding pack

TABLE XXIX: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------|----------------|----------------|
| 1 | 247.430 | 229.251 | 225.370 |
| 2 | 174.185 | 168.884 | 168.979 |
| 4 | 157.589 | 155.282 | 157.478 |
| 8 | 161.628 | 158.789 | 156.443 |
| 10 | 188.020 | 167.726 | 158.378 |
| 12 | 161.222 | 158.014 | 154.675 |
| 16 | 159.562 | 153.257 | 148.363 |

TABLE XXX: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.540 | 0.274 |
| 2 | 0.710 | 0.366 | 0.183 |
| 4 | 0.393 | 0.199 | 0.098 |
| 8 | 0.191 | 0.097 | 0.049 |
| 10 | 0.132 | 0.074 | 0.039 |
| 12 | 0.128 | 0.065 | 0.033 |
| 16 | 0.097 | 0.050 | 0.026 |

L. pytorch encoding scatter

TABLE XXXI: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------|----------------|----------------|
| 1 | 317.663 | 268.323 | 265.387 |
| 2 | 184.717 | 179.515 | 173.750 |
| 4 | 139.132 | 123.383 | 126.569 |
| 8 | 105.950 | 102.119 | 100.966 |
| 10 | 98.206 | 94.839 | 91.242 |
| 12 | 95.794 | 96.229 | 90.712 |
| 16 | 87.427 | 84.979 | 84.262 |

TABLE XXXII: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.592 | 0.299 |
| 2 | 0.860 | 0.442 | 0.229 |
| 4 | 0.571 | 0.322 | 0.157 |
| 8 | 0.375 | 0.194 | 0.098 |
| 10 | 0.323 | 0.167 | 0.087 |
| 12 | 0.276 | 0.138 | 0.073 |
| 16 | 0.227 | 0.117 | 0.059 |

M. linux encoding pack

TABLE XXXIII: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|----------------|----------------|----------------|
| 1 | 169.893 | 162.189 | 146.131 |
| 2 | 120.387 | 116.712 | 115.819 |
| 4 | 116.030 | 111.948 | 112.458 |
| 8 | 113.034 | 111.988 | 113.310 |
| 10 | 113.167 | 115.369 | 121.193 |
| 12 | 120.443 | 121.303 | 145.063 |
| 16 | 141.595 | 123.439 | 126.436 |

TABLE XXXIV: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.524 | 0.291 |
| 2 | 0.706 | 0.364 | 0.183 |
| 4 | 0.366 | 0.190 | 0.094 |
| 8 | 0.188 | 0.095 | 0.047 |
| 10 | 0.150 | 0.074 | 0.035 |
| 12 | 0.118 | 0.058 | 0.024 |
| 16 | 0.075 | 0.043 | 0.021 |

N. linux encoding scatter

TABLE XXXV: Overall encoding times

| Threads Processes | 1 | 2 | 4 |
|----------------------|---------|---------------|----------------|
| 1 | 189.447 | 150.256 | 147.334 |
| 2 | 122.175 | 118.490 | 117.440 |
| 4 | 84.408 | 81.996 | 81.776 |
| 8 | 68.110 | 67.038 | 43.153 |
| 10 | 63.032 | 61.393 | 61.122 |
| 12 | 62.554 | 61.992 | 62.346 |
| 16 | 57.539 | 56.922 | 56.938 |

TABLE XXXVI: Overall encoding efficiency

| Threads Processes | 1 | 2 | 4 |
|----------------------|-------|-------|-------|
| 1 | 1.000 | 0.630 | 0.321 |
| 2 | 0.775 | 0.400 | 0.202 |
| 4 | 0.561 | 0.289 | 0.145 |
| 8 | 0.348 | 0.177 | 0.137 |
| 10 | 0.301 | 0.154 | 0.077 |
| 12 | 0.252 | 0.127 | 0.063 |
| 16 | 0.206 | 0.104 | 0.052 |

O. Local testing

TABLE XXXVII: Overall encoding times for different folders on local machine

| Source | Processes Threads | 1 1 | 2 1 | 4 1 |
|---------------------|----------------------|---------|---------|---------|
| Numpy | | | | |
| Encoding times | | 2.0731 | 1.0857 | 0.621 |
| Encoding efficiency | | 1.0 | 0.9547 | 0.8345 |
| Decoding times | | 1.8602 | 1.0178 | 0.5757 |
| Decoding efficiency | | 1.0 | 0.9138 | 0.8078 |
| Pytorch | | | | |
| Encoding times | | 15.1102 | 7.7887 | 5.4025 |
| Encoding efficiency | | 1.0 | 0.97 | 0.6992 |
| Decoding times | | 14.6897 | 7.7248 | 5.6166 |
| Decoding efficiency | | 1.0 | 0.9508 | 0.6538 |
| Linux | | | | |
| Encoding times | | 85.0387 | 49.5736 | 30.6647 |
| Encoding efficiency | | 1.0 | 0.8577 | 0.6933 |
| Decoding times | | 80.3765 | 47.8856 | 29.6687 |
| Decoding efficiency | | 1.0 | 0.8393 | 0.6773 |